

Diwesh Sharma

 diweshsharma |  diweshsharma |  diweshsharma.janu@gmail.com |  +91-9354680059

Summary

Proficient in machine learning and NLP, building end-to-end pipelines using Python, Scikit-learn, XGBoost, and Streamlit. Experienced in data preprocessing, feature engineering, model evaluation, and deployment via FastAPI and Docker. Built and deployed projects achieving 94% accuracy in NLP classification and 0.97 ROC-AUC in anomaly detection. Strong foundation in Data Structures, Statistics, and Algorithms. Actively seeking AI/ML Engineer roles.

Education

Dronacharya College of Engineering

Bachelor of Technology, Computer Science (AI & ML) | CGPA: 8.0/10.0

A.V.R Public School

Senior Secondary (XII), Science, CBSE | 89%

A.V.R Public School

Secondary School (X), CBSE | 93%

Gurugram, Haryana

May 2023 – May 2027

Delhi

2023

Delhi

2021

Skills

Programming Languages: Python, SQL, C++

ML/AI Frameworks: Scikit-learn, XGBoost

Libraries & Tools: NumPy, Pandas, Matplotlib, Seaborn, NLTK

Data Science: NLP, Feature Engineering, Statistical Analysis, Model Evaluation, Data Visualization

Technologies: FastAPI, Streamlit, Docker, Git, GitHub, Jupyter Notebook

Projects

Student Placement Prediction System | *Python, Scikit-learn, XGBoost, FastAPI, Groq API, Render* [Demo Link](#)

- Trained and benchmarked 6 models (Logistic Regression, Random Forest, XGBoost for classification; Linear Regression, Random Forest, XGBoost Regressor for salary prediction) on 100K student records
- XGBoost selected as best classifier (Recall: 0.63) and Linear Regression as best regressor (R^2 : 0.79); deployed cascaded pipeline with shared preprocessor to eliminate data leakage
- Integrated Groq LLM (Llama 3.1) for personalized weak-area suggestions; deployed full-stack on Render with FastAPI backend and Streamlit frontend

Credit Card Fraud Detection | *Python, XGBoost, SMOTE, Scikit-learn, Imbalanced-learn*

[GitHub](#)

- Achieved 98% recall and 0.97 ROC-AUC on highly imbalanced dataset (0.17% fraud rate) using XGBoost with hyperparameter tuning
- Applied SMOTE oversampling to address class imbalance, improving fraud detection rate by 35%
- Benchmarked Logistic Regression, Decision Tree, and XGBoost; optimized for recall and F1-score to minimize false negatives

Fake News Detection System | *Python, NLP, TF-IDF, NLTK, Scikit-learn*

[GitHub](#)

- NLP classification pipeline achieving 94% accuracy on 10,000+ articles using TF-IDF vectorization and Random Forest classifier

Heart Disease Prediction | *Python, Scikit-learn, Streamlit, Plotly*

[Demo Link](#)

- Deployed Streamlit app predicting heart disease risk on UCI dataset (303 patients, 13 features) achieving 84.49% accuracy with interactive Plotly dashboard

Certifications

Certified Generative AI Professional – Oracle

Explore Machine Learning using R – Infosys